

Operation & Metric Analytics











- **Project Description:** This project aims to analyze user activity, engagement metrics, and operational data to extract valuable insights for performance optimization and strategic decision-making. Key tasks include calculating throughput, job review trends, language shares analysis, investigating metric spikes, and analyzing user growth, retention, and email engagement. The tasks were executed systematically using SQL queries to ensure data accuracy and generate actionable insights.

A) Job Data Analysis:






1. **Jobs Reviewed Over Time:** Write an SQL query to calculate the number of jobs reviewed per hour for each day in November 2020.
- **Approach:** To analyze the number of jobs reviewed per hour for each day in November 2020, the dataset is filtered so that only records within the specified date range (2020-11-01 to 2020-11-30) are included. By using SQL date and time are extracted from the ds column and then grouped by these time intervals. The query calculates the total jobs reviewed for each combination of date and hour, ensuring a chronological order for easy trend analysis.
 - **Insights:**
 - **Daily Trends:** This query shows job review activity per hour, helping to identify peak working hours and low activity periods.
 - **Operational Patterns:** Insights into hourly workloads can assist in resource allocation and process optimization.
 - **Spike Detection:** It is possible to locate specific timeframes with unusual spikes in job reviews using hourly data, which can be further investigated.

- **Result:**




Operation & Metric Analytics* x



Limit to 1000 rows



```
16 /* 1. Jobs Reviewed Over Time: Write an SQL query to calculate the r
17 in November 2020*/
18 • SELECT DATE(ds) AS review_date,
19         HOUR(ds) AS review_hour,
20         COUNT(job_id) AS jobs_reviewed
21 FROM job_data
22 WHERE
23     DATE(ds) BETWEEN '2020-11-01' AND '2020-11-30'
24 GROUP BY
25     review_date, review_hour
26 ORDER BY
27     review_date, review_hour;
28
```

Result Grid  Filter Rows: Export:  Wrap Cell Content: 

	review_date	review_hour	jobs_reviewed
▶	2020-11-25	0	1
	2020-11-26	0	1
	2020-11-27	0	1
	2020-11-28	0	2
	2020-11-29	0	1
	2020-11-30	0	2

2. Throughput Analysis: Write an SQL query to calculate the 7-day rolling average of throughput. Additionally, explain whether you prefer using the daily metric or the 7-day rolling average for throughput, and why.

- **Approach:**

I used CTE to calculate the 7-day rolling average of throughput:

- i. The **daily_throughput** CTE computes the number of events per second for each day by dividing the total daily events by the total seconds in a day ($24 * 60 * 60$).
- ii. The **rolling_avg** CTE calculates the rolling average of events per second over a 7-day window using the AVG() window function.

- **Preference:**

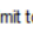












I prefer the 7-day Rolling Average for throughput analysis as it provides a balanced view by smoothing short-term variations, making it easier to identify overall trends without being overly affected by daily noise. On the other hand, the daily metric might be more appropriate if decision-making depends on immediate daily variations.

- **Insights:**






- **7-Day Rolling Average:** Offers a smoother trend by mitigating the effect of daily fluctuations and outliers, making it better for long-term monitoring.
- **Daily Metric:** Useful for real-time monitoring and detecting abrupt spikes but can be noisy.
- **Conclusion:** Prefer 7-day rolling averages for stable insights and trend analysis.

- **Result:**

Operation & Metric Analytics*



Limit to 1000 rows



```
29  /* 2. Throughput Analysis: Write an SQL query to calculate the 7-day rolling average of
30  • WITH daily_throughput AS (
31      SELECT DATE(ds) AS event_date,
32             COUNT(*) AS total_events,
33             COUNT(*) * 1.0 / (24 * 60 * 60) AS events_per_second
34      FROM job_data
35      GROUP BY DATE(ds)),
36  rolling_avg AS (
37      SELECT event_date,
38             events_per_second,
39             AVG(events_per_second)
40             OVER (ORDER BY event_date
41                  ROWS BETWEEN 6 PRECEDING AND CURRENT ROW) AS rolling_avg_throughput
42      FROM daily_throughput)
43  SELECT
44      event_date, events_per_second, rolling_avg_throughput
45  FROM rolling_avg
46  ORDER BY event_date;
```

Result Grid

Filter Rows:

Export:

Wrap Cell Content:

	event_date	events_per_second	rolling_avg_throughput
▶	2020-11-25	0.00001	0.000010000
	2020-11-26	0.00001	0.000010000
	2020-11-27	0.00001	0.000010000
	2020-11-28	0.00002	0.000012500
	2020-11-29	0.00001	0.000012000
	2020-11-30	0.00002	0.000013333

Result 44 x

Output

3. Language Share Analysis: Write an SQL query to calculate the percentage share of each language over the last 30 days.

- **Approach:**

I used CTE, to calculate the percentage share of each language in the last 30 days:

- i. Calculate the number of job days for each language in the last 30 using the recent_jobs CTE.
- ii. Calculate the overall job count using the total_jobs CTE.
- iii. Calculate each language's percentage share by dividing its job count by the total count.

- **Insights:**

- The dominant languages have been identified in the past 30 days.
- **Persian** is the dominant language with a **37.50%** language share.
- Identifies trends in language preferences or helps allocate resources.

- **Result:**

Operation & Metric Analytics* x

```

47  /* 3. Language Share Analysis: Write an SQL query to calculate the percentage share of each language
48  • WITH recent_jobs AS (
49      SELECT language,
50             COUNT(*) AS job_count
51      FROM job_data
52      WHERE DATE(ds) >= DATE_SUB('2020-12-01', INTERVAL 30 DAY)
53      GROUP BY language
54  ),
55  total_jobs AS (
56      SELECT SUM(job_count) AS total_job_count
57      FROM recent_jobs
58  )
59  SELECT
60      language, job_count,
61      ROUND((job_count * 100.0) / total_job_count, 2) AS language_share_percentage
62  FROM recent_jobs
63  CROSS JOIN total_jobs
64  ORDER BY language_share_percentage DESC;
  
```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: [\[A\]](#)

	language	job_count	language_share_percentage
▶	Persian	3	37.50
	English	1	12.50
	Arabic	1	12.50
	Hindi	1	12.50
	French	1	12.50
	Italian	1	12.50

Result 13 x

4. Duplicate Rows Detection: Write an SQL query to display duplicate rows from the job_data table.

- **Approach:**

- To display duplicate rows from the job_data table, I used queries to identify exact duplicates across all columns and partial duplicates.
- For exact duplicates, I used the GROUP BY clause for all columns and filtered using HAVING COUNT(*) > 1.
- For partial duplicates based on certain columns (e.g. job_id and actor_id), I used the GROUP BY clause for specific columns and filtered similarly.

- **Insights:**

- There are no rows where all column values are identical, which means the dataset not contain any full duplicates.
- No duplicates rows are returned, the dataset does not have exact duplicates.
- **job_id 23** appears multiple times on different **dates** but with different **actor_id** or **event**.
- **actor_id 1003** is also repeated on different **dates** with different **job_id** or **language**.

- **Result:**

[illegible]

B) Investigating Metric Spike:

1. Weekly User Engagement: Write an SQL query to calculate the weekly user engagement.

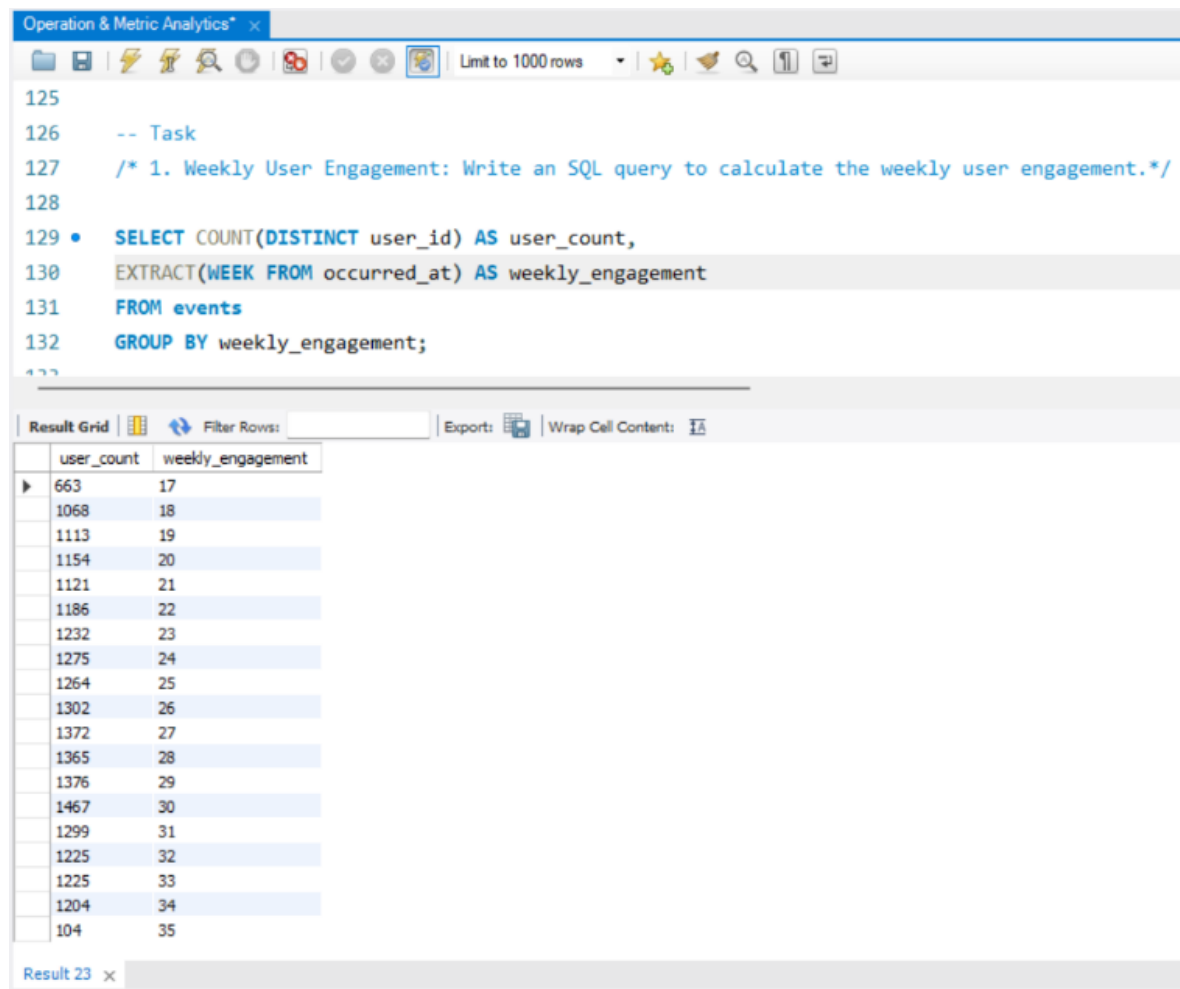
- **Approach:**

- Calculate the number of unique users (COUNT(DISTINCT user_id)) for each week (EXTRACT(WEEK FROM occurred_at)).
- Group results by week to track weekly user engagement.

- **Insights:**

Helps to monitor the number of active users each week and identify trends or anomalies.

- **Result:**



The screenshot shows a SQL IDE window titled "Operation & Metric Analytics". The query editor contains the following SQL code:

```
-- Task
/* 1. Weekly User Engagement: Write an SQL query to calculate the weekly user engagement.*/

SELECT COUNT(DISTINCT user_id) AS user_count,
       EXTRACT(WEEK FROM occurred_at) AS weekly_engagement
FROM events
GROUP BY weekly_engagement;
```

Below the query editor, the "Result Grid" tab is active, displaying the results of the query. The results are shown in a table with two columns: "user_count" and "weekly_engagement". The table contains 17 rows of data, showing a steady increase in user count over time.

	user_count	weekly_engagement
▶	663	17
	1068	18
	1113	19
	1154	20
	1121	21
	1186	22
	1232	23
	1275	24
	1264	25
	1302	26
	1372	27
	1365	28
	1376	29
	1467	30
	1299	31
	1225	32
	1225	33
	1204	34
	104	35

At the bottom of the window, the status bar indicates "Result 23 x".

2. User Growth Analysis: Write an SQL query to calculate the user growth for the product.

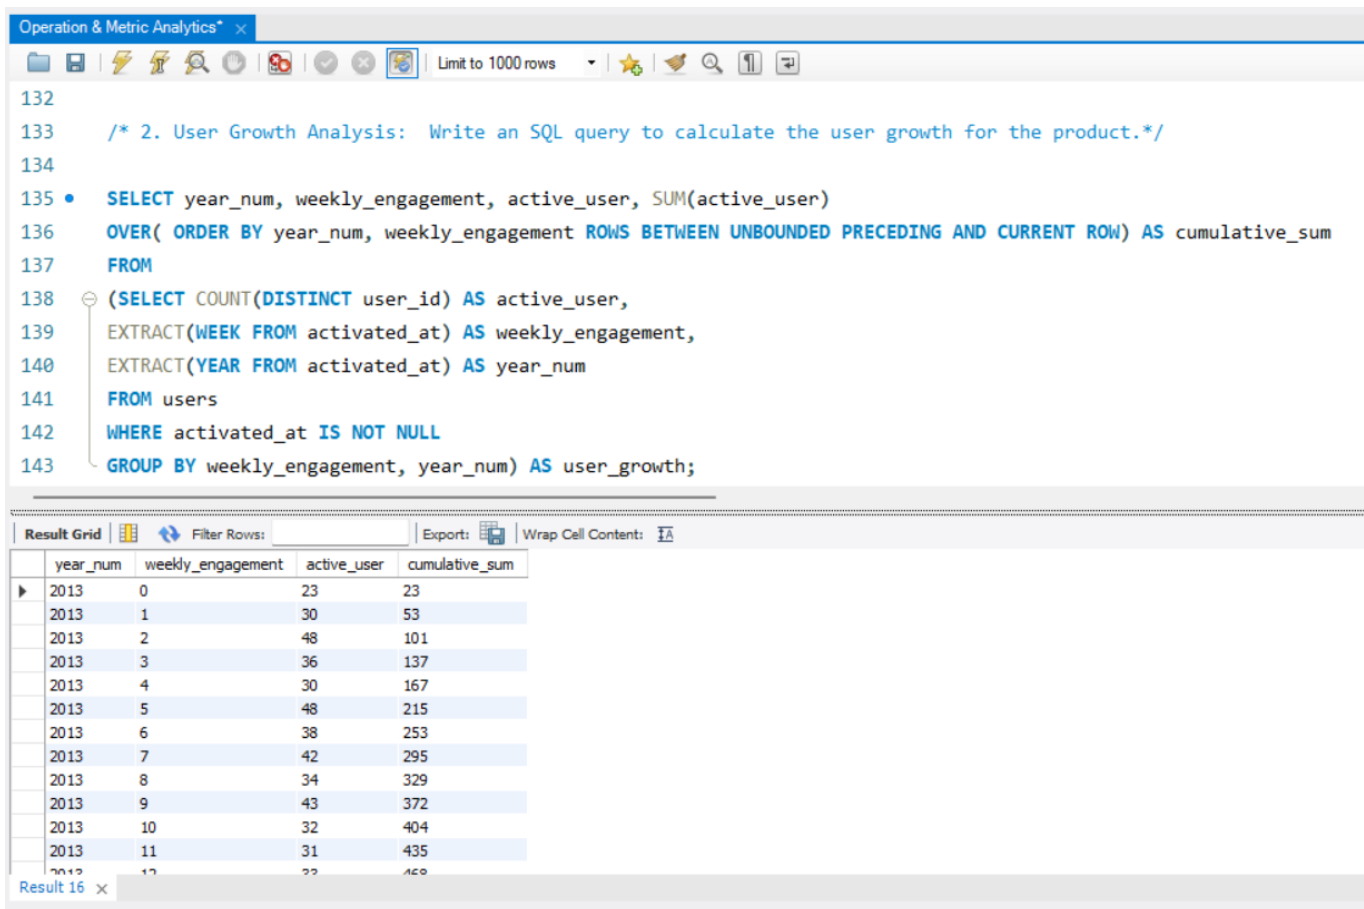
- **Approach:**

- Compute the weekly number of active users by counting distinct user IDs (COUNT(DISTINCT user_id)) grouped by week and year of activation.
- Use a window function (SUM() with OVER) to calculate the cumulative sum of active users to track overall growth.

- **Insights:**

- Tracks user acquisition over time and reveals growth patterns.
- Helps assess the success of marketing or onboarding efforts.

- **Result:**



The screenshot shows a SQL IDE window titled "Operation & Metric Analytics". The query editor contains the following SQL code:

```
/* 2. User Growth Analysis: Write an SQL query to calculate the user growth for the product.*/  
  
SELECT year_num, weekly_engagement, active_user, SUM(active_user)  
OVER( ORDER BY year_num, weekly_engagement ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) AS cumulative_sum  
FROM  
(SELECT COUNT(DISTINCT user_id) AS active_user,  
EXTRACT(WEEK FROM activated_at) AS weekly_engagement,  
EXTRACT(YEAR FROM activated_at) AS year_num  
FROM users  
WHERE activated_at IS NOT NULL  
GROUP BY weekly_engagement, year_num) AS user_growth;
```

The results are displayed in a table with the following columns: year_num, weekly_engagement, active_user, and cumulative_sum. The data shows weekly growth from 2013, with the cumulative sum increasing over time.

year_num	weekly_engagement	active_user	cumulative_sum
2013	0	23	23
2013	1	30	53
2013	2	48	101
2013	3	36	137
2013	4	30	167
2013	5	48	215
2013	6	38	253
2013	7	42	295
2013	8	34	329
2013	9	43	372
2013	10	32	404
2013	11	31	435
2013	12	22	457

3. Weekly Retention Analysis: Write an SQL query to calculate the weekly retention of users based on their sign-up cohort.

- **Approach:**

- Identify a cohort of users based on their signup week (cohort_year and cohort_week).
- Count the number of retained users (COUNT(DISTINCT wa.user_id)) for subsequent weeks by combining this cohort data with user activity.
- Calculate the week number difference between signup and activity weeks.

- **Insights:**

- Analyzes retention trends by week, highlighting how well users are retained over time.
- This is useful for identifying issues with product engagement after signing up.

- **Result:**

```
Operation & Metric Analytics x
Limit to 1000 rows
145
146 /* 3. Weekly Retention Analysis: Write an SQL query to calculate the weekly retention of users based on their
147 sign-up cohort.*/
148
149 WITH cohort AS (
150     SELECT user_id,
151            EXTRACT(YEAR FROM created_at) AS cohort_year,
152            EXTRACT(WEEK FROM created_at) AS cohort_week
153     FROM users),
154 weekly_activity AS (
155     SELECT user_id,
156            EXTRACT(YEAR FROM occurred_at) AS activity_year,
157            EXTRACT(WEEK FROM occurred_at) AS activity_week
158     FROM events),
159 retention AS (
160     SELECT c.cohort_year, c.cohort_week,
161            wa.activity_year, wa.activity_week,
162            COUNT(DISTINCT wa.user_id) AS retained_users
163     FROM cohort c
```

Operation & Metric Analytics* x

Limit to 1000 rows

```

163 FROM cohort c
164 JOIN weekly_activity wa
165 ON c.user_id = wa.user_id
166 GROUP BY c.cohort_year, c.cohort_week, wa.activity_year, wa.activity_week),
167 weekly_retention AS (
168 SELECT
169     cohort_year, cohort_week,
170     activity_year, activity_week,
171     retained_users,
172     activity_week - cohort_week AS week_number
173 FROM retention
174 WHERE
175     activity_year = cohort_year
176     AND activity_week >= cohort_week)
177 SELECT
178     cohort_year, cohort_week,
179     week_number, retained_users
180 FROM weekly_retention
181 ORDER BY cohort_year, cohort_week, week_number;

```

Result Grid | Filter Rows: | Export: | Wrap Cell Co

	cohort_year	cohort_week	week_number	retained_users
▶	2014	0	17	3
	2014	0	18	8
	2014	0	19	12
	2014	0	20	9
	2014	0	21	10
	2014	0	22	10
	2014	0	23	9
	2014	0	24	13
	2014	0	25	8
	2014	0	26	10
	2014	0	27	8
	2014	0	28	7
	2014	0	29	6
	2014	0	30	7
	2014	0	31	6
	2014	0	32	5
	2014	0	33	4
	2014	0	34	3
	2014	1	16	4
	2014	1	17	10
	2014	1	18	12
	2014	1	19	12
	2014	1	20	16
	2014	1	21	8
	2014	1	22	11
	2014	1	23	11

Result 32 x

4. Weekly Engagement Per Device: Write an SQL query to calculate the weekly engagement per device.

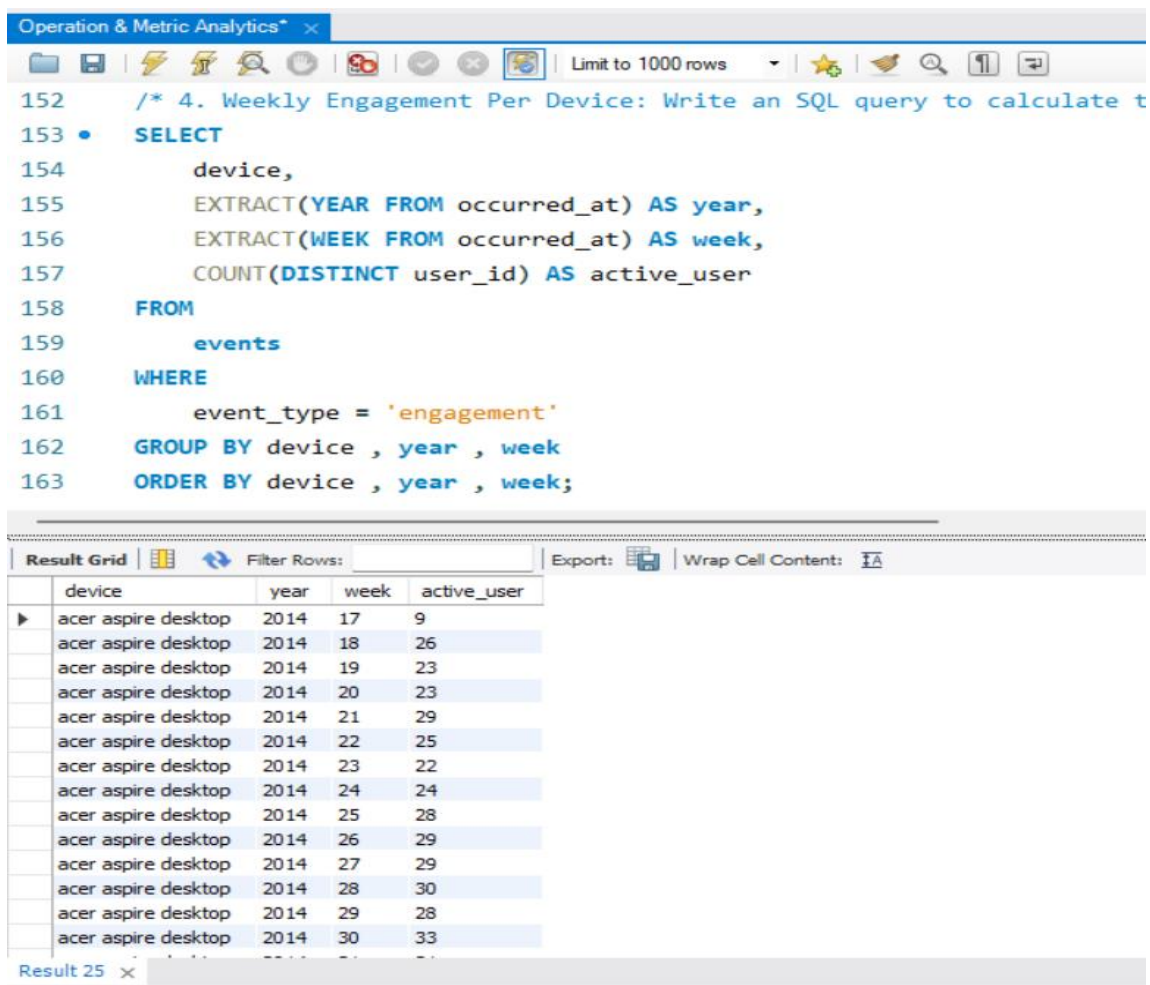
- **Approach:**

- Grouping user activity by device, year, and week to count the number of distinct active users (COUNT(DISTINCT user_id)).
- Filtered data where the event type is engagement.

- **Insights:**

- Identifies device-specific engagement trends.
- Helps optimize for devices with higher engagement rates.

- **Result:**



The screenshot shows a SQL IDE window titled "Operation & Metric Analytics* x". The query editor contains the following SQL query:

```
152  /* 4. Weekly Engagement Per Device: Write an SQL query to calculate t
153  • SELECT
154      device,
155      EXTRACT(YEAR FROM occurred_at) AS year,
156      EXTRACT(WEEK FROM occurred_at) AS week,
157      COUNT(DISTINCT user_id) AS active_user
158  FROM
159      events
160  WHERE
161      event_type = 'engagement'
162  GROUP BY device , year , week
163  ORDER BY device , year , week;
```

Below the query editor, the "Result Grid" tab is active, displaying the results of the query. The table has four columns: device, year, week, and active_user. The results show engagement data for "acer aspire desktop" from 2014, week 17 to 30.

device	year	week	active_user
acer aspire desktop	2014	17	9
acer aspire desktop	2014	18	26
acer aspire desktop	2014	19	23
acer aspire desktop	2014	20	23
acer aspire desktop	2014	21	29
acer aspire desktop	2014	22	25
acer aspire desktop	2014	23	22
acer aspire desktop	2014	24	24
acer aspire desktop	2014	25	28
acer aspire desktop	2014	26	29
acer aspire desktop	2014	27	29
acer aspire desktop	2014	28	30
acer aspire desktop	2014	29	28
acer aspire desktop	2014	30	33

The bottom of the window shows "Result 25 x".

5. Email Engagement Analysis: Write an SQL query to calculate the email engagement metrics.

- **Approach:**

Email engagement metrics are data points that provide information about how users interact with your emails such as Click-through rate (CTR) and Open rate. For calculating email engagement metrics.

1) Email Events Count:

- a. Count the total number of events grouped by action.

2) Engagement Rates:

- a. Classify events into categories (**email_sent, email_open, email_click**) using a **CASE statement**.
- b. Calculate open and click rates as percentages based on the number of emails sent.

- **Insights:**

- Tracks email effectiveness through open and click rates.
- Helps refine email campaigns by identifying which emails perform well.

- **Result:**

```
Operation & Metric Analytics ×
Limit to 1000 rows
172
173 • SELECT
174 100.0 * SUM(CASE WHEN email_category='email_open' then 1 else 0 end)
175 / SUM(CASE WHEN email_category='email_sent' then 1 else 0 end) AS open_rate,
176 100.0 * SUM(CASE WHEN email_category='email_click' then 1 else 0 end)
177 / SUM(CASE WHEN email_category='email_sent' then 1 else 0 end) AS click_rate
178 FROM
179 (
180 SELECT *,
181 CASE
182 WHEN action IN('sent_weekly_digest', 'sent_reengagement_email') THEN 'email_sent'
183 WHEN action IN('email_open') THEN 'email_open'
184 WHEN action IN('email_clickthrough') THEN 'email_click'
185 END AS email_category
186 FROM email_events
187 ) AS email_metrics;
188
189
```

	open_rate	click_rate
▶	33.58339	14.78989

- **Tech-Stack Used**

MySQL Workbench (Version 8.0): Used for writing and executing SQL queries due to its user-friendly interface and ability to visualize results.

- **Drive Link**

SQL File:

https://drive.google.com/file/d/1K0oGjcYRFgXAhWkBL_dMEU2xqMQ7hitu/view?usp=sharing

Project Link:

https://drive.google.com/drive/folders/1DxveRXxukxHPx2Vby1eIH0CiPUfRT_N?usp=sharing